

Distributed PageRank - Final Report

Aaron Myers, Megan Ruthven

December 16, 2014

Contents

1	Introduction	2
2	Related Work	2
3	Load Balancing	3
4	Algorithms and Implementation	3
4.1	Topology-driven	3
4.2	Data-Driven Pagerank	3
4.3	Linear System Approach	7
4.4	ADMM	8
5	Datasets	10
6	Results	10
6.1	Power Iteration Results	10
6.2	ADMM Results Compared to other Linear methods	13
7	Future Work	15
8	Conclusion	15
	Appendices	17

1 Introduction

The purpose of this project is to investigate distributed PageRank by applying Pagerank methods currently in use in the multi-thread case and implement it in the distributed setting. It is to determine if these methods can be scaled to multi-machine systems. We also take a method that has already proven to be effective in the distributed setting and apply it to the Pagerank problem Alternating Direction Method of Multipliers (ADMM) [1]). Immediately below is a more detailed description of each approach. The primary metrics for performance will be: speedup, scalability, and ease of implementation, all of which will be explained in the results section. The last metric is included to suggest that ease of coding and understanding of a method has an impact on the adoption rate in industry and therefore should be included in this investigation.

We chose to investigate ADMM [1] because it would have higher performance for ease of coding and understanding. We applied ADMM to the linear form of the Pagerank problem ($Ax=b$), but it is more difficult to apply a data-driven approach to the separated minimization problems, so to investigate data-driven approaches outlined in [2], we implemented data-driven power iteration with MPI and openMP. Typical power iteration will be referred to as Topology driven. The data-driven methods were Pull, Pull-Push, and Push based. The details of each permutation of the data-driven methods will be discussed in following sections. Because each method is implemented on distributed nodes on a big dataset, each method took advantage of the separate memories, and divided the connection matrices between nodes. There, we implemented a static load balancing regime based on amount of nonzeros in the assigned rows, which proved to be more efficient than load balancing on the number of x elements to calculate. Load balancing will be explained more in the following sections.

The methods' performance will be compared on total time, speed up, and scalability over varying amount of nodes (MPI), and cores per node (openMP). And finally, we will go over the limitations of the methods in MPI and openMP, and future work.

2 Related Work

Pagerank is the algorithm that Larry Page and Sergey Brin researched, and subsequently used their research to form the company Google. Since then, many researchers have focused on improving different aspects of Pagerank. Our project focused on different implementations of Pagerank to compare and contrast implementations' performance in a distributed setting.

Iterative solvers for the formulation $Ax=b$ using well known methods including but not limited to: GMRES, BiCGSTAB, CGS, Chebychev iterations, along with Jacobi iteration methods [3]. Their results indicated that normal Power Iteration and Jacobi methods almost always performed better than the other linear solvers and although we do not have access to their data sets or machines, we will attempt to generally compare our speedup curves to theirs and expect some correlation. We will also expect that our power iteration and data-driven approaches will perform better than our ADMM solver for $Ax=b$ as they did for [3].

Previous implementations of Pagerank with MPI [4] showed that it is expected for compute time to decrease as the number of MPI processes increase to a certain point. As seen in the figure presented in the paper, with a 1M node graph, after about 30 cores, the solve time starting to trend slightly upward. We assume this increase in solve time is primarily due to the increase in message passing overhead. This would suggest that we should see similar results and also that there may be an optimal core number depending on the size of the data.

We have worked extremely closely with Joyce Whang as many of our algorithms were taken

directly from her research paper and other members of her research group [2]. In her submitted paper, the algorithms proposed are implemented on a Galois multithread system. Our data-driven implementations are taken directly from her paper and implemented with the MPI and openMP frameworks. This extraction to the distributed setting is to confirm or deny the idea that her algorithms could also be effective across nodes.

3 Load Balancing

Before this paper details the different algorithms we implemented, it will talk about load balancing. Load balancing on distributed systems could be scheduled in task or job manner, as discussed in [5], where jobs were a collection of tasks. Because MPI is best suited for batch passing of messages, our project used a job based approach. The two approaches were to balance jobs by number of tasks (indices to compute Pagerank), or balance by total average work for each node (summation of nonzeros of each incoming node for the assigned indices). First we implemented our methods where each node had the same amount of tasks to compute. Then, we implemented balancing by total work, and compared the differences. The topological and data-driven Pagerank algorithms consistently sped up where the total work balancing was about two times faster than the balanced tasked method. For these reasons, the results reflect algorithms which used the work balancing method.

4 Algorithms and Implementation

This section discusses the implementations of different methods for computing pagerank. This includes topological, data-driven, and ADMM Pagerank.

4.1 Topology-driven

This is the classic implementation for calculating Pagerank, seen in algorithm 1. It is straight forward, but recalculates all X values every iteration. This could be redundant on a large percentage of the X indices, and therefore, it could waste computational power.

In the topological algorithm, every x_v is updated. Each node computes a subsection of \mathbf{x} , named \mathbf{x}_r , then all of the nodes combine their new result in batch to create a new \mathbf{x} in order to continue to compute the newer values of \mathbf{x}_r . All of the recomputation within each node only updates the current x_v , this creates for a thread safe algorithm.

4.2 Data-Driven Pagerank

In addition to the ADMM implementation, we implemented power iteration, as well as three permutations of the data-driven pagerank method (pull, pull-push, and push). The data-driven method (taken from [2]) aims to minimize unnecessary computation by only updating pageranks of elements whose incoming connections were updated to a satisfactory degree.

The Pull data-driven method in algorithm 2 selectively updates the values of x_v based on if an incoming connected node updated itself above the threshold, ϵ . This is determined by the array t_c , which is a set of c indices for that node. In the computation of x_v , if the difference between the new and old value is above ϵ , set all outgoing nodes' t_w to true. Notice that each element, v , updates its value, x_v , and indices of t both in and out of set c . This requires the syncing of all copies of t^{new} between nodes with a logical or. Additionally, this algorithm accesses more memory by accessing incoming nodes' pagerank and outgoing nodes' t^{new} .

Algorithm 1 Topology-driven Pagerank

```
1: Input: graph  $P_r = (V_r, E_r)$ ,  $\alpha$ ,  $\epsilon$ 
2: Output: Pagerank  $\mathbf{x}$ 
3: Initialize  $\mathbf{x} = \alpha \mathbf{e}$ 
4:  $\exists \mathbf{x}_r$  in  $\mathbf{x}$  as  $P_r$  is to  $P$ 
5: while true do
6:   for  $v \in V_r$  do
7:      $x_v^{(k+1)} = \alpha + (1 - \alpha) \sum_{w \in S_v} \frac{x_w^{(k)}}{|T_w|}$ 
8:      $\delta_v = |x_v^{(k+1)} - x_v^{(k)}|$ 
9:   end for
10:  sync all  $\mathbf{x}_r$  between nodes
11:  if  $\|\delta\|_\infty < \epsilon$  then
12:    break;
13:  end if
14: end while
15:  $\mathbf{x} = \frac{\mathbf{x}}{\|\mathbf{x}\|_1}$ 
```

Algorithm 2 Pull Data-driven Pagerank

```
1: Input: graph  $P_c = (V_c, E_c)$ ,  $\alpha$ ,  $\epsilon$ 
2: Output: Pagerank  $\mathbf{x}$ 
3: Initialize  $\mathbf{x} = \alpha \mathbf{e}$  and  $\mathbf{t}_c = true \times \mathbf{e}_c$ 
4:  $\exists \mathbf{x}_c$  in  $\mathbf{x}$  as  $P_c$  is to  $P$ 
5: while  $\exists v$  s.t.  $t_v = true$  do
6:    $t^{new} = false \times \mathbf{e}$ 
7:   for  $v \in V_c$  do
8:     if  $t_v = true$  then
9:        $x_v^{new} = \alpha + (1 - \alpha) \sum_{w \in S_v} \frac{x_w}{|T_w|}$ 
10:      if  $|x_v^{new} - x_v| \geq \epsilon$  then
11:         $x_v = x_v^{new}$ 
12:        for  $w \in T_v$  do
13:           $t_w^{new} = true$ 
14:        end for
15:      end if
16:    end if
17:  end for
18:  sync all  $\mathbf{x}_c$  between nodes
19:  logical or all  $t = t^{new}$  between nodes
20: end while
21:  $\mathbf{x} = \frac{\mathbf{x}}{\|\mathbf{x}\|_1}$ 
```

Algorithm 3 Pull-Push Data-driven Pagerank

```
1: Input: graph  $P_c = (V_c, E_c)$ ,  $\alpha$ ,  $\epsilon$ 
2: Output: Pagerank  $\mathbf{x}$ 
3: Initialize  $\mathbf{x} = \alpha \mathbf{e}$ 
4:  $\exists \mathbf{x}_c$  in  $\mathbf{x}$  as  $P_c$  is to  $P$ 
5: Initialize  $\mathbf{r} = \mathbf{0}$ 
6: for  $v \in V_c$  do
7:   for  $w \in S_v$  do
8:      $r_v = r_v + \frac{1}{|T_w|}$ 
9:   end for
10:   $r_v = (1 - \alpha)\alpha r_v$ 
11: end for
12: while  $\exists v$  s.t.  $r_v \geq \epsilon$  do
13:   $\mathbf{r}^{new} = \mathbf{0}$ 
14:  for  $v \in V_c$  do
15:    if  $r_v \geq \epsilon$  then
16:       $x_v = \alpha + (1 - \alpha) \sum_{w \in S_v} \frac{x_w}{|T_w|}$ 
17:      for  $w \in T_v$  do
18:         $r_w^{new} = r_w^{new} + \frac{r_v \alpha}{|T_v|}$ 
19:      end for
20:    else
21:       $r_v^{new} = r_v^{new} + r_v$ 
22:    end if
23:  end for
24:  sync all  $\mathbf{x}_c$  between nodes
25:  add all and scatter  $r_c = r^{new}$  between nodes
26: end while
27:  $\mathbf{x} = \frac{\mathbf{x}}{\|\mathbf{x}\|_1}$ 
```

The Pull-Push data-driven method in algorithm 3 selectively updates the values of x_v based on if its residual is above a threshold, ϵ . The residuals are accounted for in the \mathbf{r} vector and batch updated to transfer residual values between nodes. Notice that each element, v , updates its value, x_v , and indices of r both in and out of set c . This requires the syncing of all copies of r^{new} between nodes by summing. Additionally, this algorithm accesses more memory by accessing incoming nodes' pagerank and outgoing nodes' r^{new} .

Algorithm 4 Push Data-driven Pagerank

```

1: Input: graph  $P_c = (V_c, E_c)$ ,  $\alpha$ ,  $\epsilon$ 
2: Output: Pagerank  $\mathbf{x}$ 
3: Initialize  $\mathbf{x} = \alpha \mathbf{e}$ 
4:  $\exists \mathbf{x}_c$  in  $\mathbf{x}$  as  $P_c$  is to  $P$ 
5: Initialize  $\mathbf{r} = \mathbf{0}$ 
6: for  $v \in V_c$  do
7:   for  $w \in S_v$  do
8:      $r_v = r_v + \frac{1}{|T_w|}$ 
9:   end for
10:   $r_v = (1 - \alpha)\alpha r_v$ 
11: end for
12: while  $\exists v$  s.t.  $r_v \geq \epsilon$  do
13:   $r^{new} = \mathbf{0}$ 
14:  for  $v \in V_c$  do
15:    if  $r_v \geq \epsilon$  then
16:       $x_v = x_v + r_v$ 
17:      for  $w \in T_v$  do
18:         $r_w^{new} = r_w^{new} + \frac{r_v \alpha}{|T_v|}$ 
19:      end for
20:    else
21:       $r_v^{new} = r_v^{new} + r_v$ 
22:    end if
23:  end for
24:  add all and scatter  $r_c = r^{new}$  between nodes
25: end while
26: sync all  $\mathbf{x}_c$  between nodes
27:  $\mathbf{x} = \frac{\mathbf{x}}{\|\mathbf{x}\|_1}$ 

```

The Push data-driven method in algorithm 4 selectively updates the values of x_v based on if its residual is above a threshold, ϵ . It updates x_v by only using r_v . This accesses less memory by bypassing fetching all of the incoming nodes current values x_w . The residuals are accounted for in the \mathbf{r} vector and batch updated to transfer residual values between nodes. Notice that each element, v , updates its value, x_v , and indices of r both in and out of set c . This requires the syncing of all copies of r^{new} between nodes by summing.

Although MPI allowed for a distributed Pagerank algorithm, there were a couple of restrictions on the implementation because of the MPI and openMP framework we chose to use. When syncing values between nodes, MPI requires all nodes to simultaneously request and send the information they require. This means that all of the independent nodes have to update the values in batch. The original implementation in Galois [2] utilized the asynchronicity of value updates unique to

Galois to be able to have threadsafe updates, and shared residual values between threads. This is not a feature in openMP or MPI. In openMP, each thread kept track of its own residuals for the whole set of indices, and reduced after all were calculated within each node, then shared its values to sum up all of the nodes' residual values. These 2 steps of synchronization were not necessary in the Galois implementation. Additionally, the batch update of residuals and the boolean update vector meant that a meaningful worklist was impossible to keep without the vectors from the other nodes. This meant, to selectively compute a new value the Pagerank, each node had to go through each value of the residual or update vectors to check if it met the standards of update.

The new vectors of information about each index needed to be passed between nodes. More information to send corresponds with a longer amount of time spend on sending data. In the case of algorithms 2 and 3, they both send \mathbf{x}_c and an update vector. This does not scale well for big

4.3 Linear System Approach

This approach requires that we form the Pagerank problem into a different linear system ($Ax=b$) where we fundamentally look at methods of iterating or directly solving for an inverse to solve for x which would provide the list of Pagerank values. Below is a simple derivation taken from [3].

$$P' = P + dv^T \quad (1)$$

$$P'' = cP' + (1 - c)ev^T \quad (2)$$

$$x^{k+1} = P''^T x^k \quad (3)$$

Where P' and P'' are the modified PageRank matrices that have the modifications necessary to create a connected graph and add a personalization factor and e is a vector of all 1's, resulting in equation 3, the Power Iteration approach to Pagerank.

Given the additional information below, we can derive the linear system for Pagerank.

$$e^T x = x^T e = \|x\|_1 = \|x\| \quad (4)$$

$$d^T x = \|x\| - \|P^T x\| \quad (5)$$

$$x = [cP^T + c(vd^T) + (1 - c)ve^T]x \quad (6)$$

Combining the information above, we arrive at the following equation:

$$(I - cP^T)x = kv \quad (7)$$

We now have Pagerank in a linear form ($Ax=b$), where $A = I - cP^T$ and $kv = b$. If in addition, we normalize x , we have the following:

$$k = \|x\| - c\|P^T x\| = (1 - c)\|x\| + d^T x \quad (8)$$

$$k = 1 - c \quad (9)$$

4.4 ADMM

Many of the articles we encountered for solving parallel pagerank in the linear form used Jacobi iteration or some Krylov Subspace method (GMRES, BiCGSTAB, etc), but we attempted to implement something we were introduced to in this course, namely ADMM [1]. This is an extremely simple way to parallelize a linear solve. This process attempts to split the linear problem into subsections, solve separately, and combine the information in a very specific way. We will compare these results to the implementation of GMRES (Generalized Minimal RESidual method) and BiCGSTAB (Bi-Conjugate Gradient method with STABilization) for the same problem parallelizing using PETSc (a software with great tools for parallelizing linear solvers) [?]. We expect ADMM to have worse performance, measured by speedup, but we would like to quantify the loss in accuracy/time relative to the ease of implementation and scalability.

Below is a brief description of the ADMM idea and algorithm [1]. We take the linear problem and split up the data accordingly:

$$A = [A_1 \dots A_n]' \quad (10)$$

$$b = [b_1 \dots b_n]' \quad (11)$$

Our original minimization of $Ax=b$ with a certain norm and regularization on x now becomes:

$$\text{minimize} \quad \sum_{i=1}^N l_i(A_i x_i - b_i) + r(z) \quad (12)$$

$$\text{subject to} \quad x_i - z = 0 \quad \forall i \quad (13)$$

Where x_i are local variables that we force to match the global solution z at each step and N is the number of processes used to solve the problem. This method also includes an 'augmented lagrangian' term. This term is used to bring robustness to the dual ascent problem and result in convergence without criteria like strict convexity or finiteness of the function. This is discussed in much greater detail in [1]. Below is the augmented lagrangian which is used to derive the resulting algorithm.

$$L_\rho(x, y) = f(x) + y^T(Ax - b) + \frac{\rho}{2} \|Ax - b\|_2^2 \quad (14)$$

$$(15)$$

The resulting algorithm, using the augmented lagrangian presented in the ADMM method [1], is as follows:

Algorithm 5 ADMM Iteration

- 1: $x_i^{k+1} = \text{argmin}_x \quad l_i(A_i x_i - b_i) + \frac{\rho}{2} \|x_i^k - z^k - u_i^k\|_2^2$
 - 2: $z^{k+1} = \text{argmin}_z \quad r(z) + \frac{N\rho}{x} \|z^k - \bar{x}^{k+1} - \bar{u}^k\|_2^2$
 - 3: $u_i^{k+1} = u_i^k + x_i^{k+1} - z^{k+1}$
-

Where $u_i^k = \frac{1}{\rho}y_i^k$ and for our implementation, we chose the L^1 regularization term (also known as lasso) with a gradient descent solver and we also attempted a direct solve with Eigen for the minimization for x, although the results were poor. Using the lasso regularization the z update becomes soft-thresholding update. Considering the Lasso method, the updated algorithm is below.

Algorithm 6 ADMM Iteration with Lasso

- 1: $x_i^{k+1} = \operatorname{argmin}_x \|(A_i x_i - b_i)\|_2^2 + \frac{\rho}{2}\|x_i^k - z^k - u_i^k\|_2^2$
 - 2: $z^{k+1} = S_{\lambda/\rho N}(\bar{x}^{k+1} - \bar{u}^k)$
 - 3: $u_i^{k+1} = u_i^k + x_i^{k+1} - z^{k+1}$
-

Where S is defined component-wise in the following way:

$$S_{\lambda/\rho N}(x_i) := (x_i - \frac{\lambda}{\rho N})_+ - (-x_i - \frac{\lambda}{\rho N})_+ \quad (16)$$

The first approach we took to solving the minimization problem on each separate process was to use the direct inversion with Eigen. Below is the resulting equation required to be solved for direct inversion.

$$x^{k+1} := (A_i^T A_i + \rho I)^{-1}(A_i^T b_i + \rho(z^k - u_i^k)) \quad (17)$$

The above equation makes it clear that using this direct solve method would require the formation of an mxm matrix (the same size as the original matrix) which would then be dense. This renders the direct solve method nearly useless and indeed when attempting this method, the algorithm (on a small matrix, m=1770961) took 2 hours to complete 1 iteration.

Another method we investigated is the gradient descent algorithm. Below is a short description of gradient descent for those unfamiliar with the method.

Algorithm 7 Gradient Descent Algorithm

- 1: Set Threshold
 - 2: Initialize x_{n+1}, x_n s.t. $\|x_{n+1} - x_n\| > threshold$
 - 3: **while** $\|x_{n+1} - x_n\| > threshold$ **do**
 - 4: $x_{n+1} = x_n - \alpha \nabla F(x_n)$
 - 5: select α s.t. $\alpha = \min \nabla F(x_{n+1})$
 - 6: **end while**
-

Step 5 above is done with line search looking at a set of discrete values between 0.001 and 10.

5 Datasets

In order to compare performances of the algorithms in different conditions, five datasets were used. They vary in size, sparsity, and quality of symmetry. These qualities are displayed in table 5. All of these qualities influence the performance of the algorithms. The PLD, Twitter RV, and SD1 datasets were the same ones used in [2]. They are good examples of fairly sizeable datasets that

Name	Size	Indices	Symmetric
Livejournal	1.1 GB	1,770,961	Yes
Friendster	7 GB	15,853,098	Yes
PLD	10.4 GB	39,497,202	No
Twitter RV	23.5 GB	41,652,230	No
SD1	32.7 GB	82,924,685	No

6 Results

The different methods had varying quantitative and qualitative results. These included the time to complete, speedup, scalability, and ease of coding.

6.1 Power Iteration Results

The Topology and Data-driven Pagerank implementations were written in C++ and used MPI to run on distributed nodes, and openMP to run multiple threads within each node. Each algorithms performance was timed on each of the datasets in all of the combinations of 1, 4, 8, 16, 32 compute nodes and 1, 4, 8, and 16 threads. That means that each algorithm for each dataset ranges from 1 total thread to 512 threads. This large scale parallelizable capacity is the utility for distributed systems. Ideally, at least one of our proposed implementations would continue to decrease it's run time as more nodes and threads are added.

Coding the different methods had many factors to consider to get the best run time from each of the methods. The code had to consider:

- Data structures to use for memory speed access and MPI sendability
- Message passing schemes to minimize message overhead time
- Thread safety
- Load balancing
- Minimizing computation
- Accessing temporal and spatial location of memory accesses
- Physical limitations of the openMP and MPI frameworks

Because of that, we believe that there are a lot of ways to minimize the computation time of the proposed methods. The solutions presented could be at a local minimum and more time spent looking at the specific implementation could unveil a implementation with lower run time. With that being said, our implementation has already undergone extensive amounts of optimization, and we feel that all of the implementations are at a level we can compare them at.

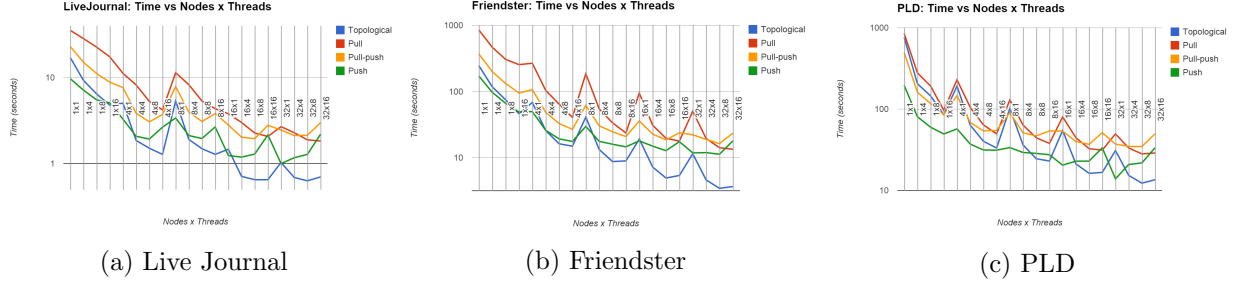


Figure 1: Convergence times of different PageRank implementations on datasets from table 5




Figure 2: Scalability over openMP threads of different PageRank implementations on datasets from table 5

For additional understanding of how performance changed over number of nodes and threads, scalability and speedup are visualized in figures 1-5. These measures are the same measures used in [2].


$$Scalability = \frac{run - time\ of\ method\ m\ with\ a\ single\ thread}{run - time\ of\ method\ m\ with\ t\ threads} \quad (18)$$

$$Speedup = \frac{run - time\ of\ the\ fastest\ single - thread\ method}{run - time\ of\ method\ m\ with\ t\ threads} \quad (19)$$


Scalability, equation 18, is a metric to show the change in runtime from one methods baseline, whereas speedup, equation 19, is to show the change in runtime from the best case scenario method at baseline configuration. Both frameworks have different artifacts that enable or hinder speedup and scalability as the number of threads increase. So, speedup and scalability are evaluated for MPI and openMP separately.

Please note on figures 2 and 4, the metric is reset at every thread = 1. This creates a misleading dip in the curve from 16 to 1 threads. Additionally, figures for the datasets not shown here are in the appendix at figures 6-10. Note Twitter RV and SD1 datasets are larger, as shown in table 5. And, because of the difference in performance of the Twitter RV and SD1, it is promising for the performance of other big datasets.


The general trends in performance as seen in 1, are that the Pull and Pull-push implementations are generally the slowest, and the Push and Topological implementations change as #1 and #2 in speed over nodes x threads. The Push algorithm outperforms the Topological method generally



(a) Live Journal




(b) Friendster




(c) PLD


Figure 3: Scalability over MPI nodes of different PageRank implementations on datasets from table 5



(a) Live Journal




(b) Friendster




(c) PLD


Figure 4: Speedup over openMP threads of different PageRank implementations on datasets from table 5



(a) Live Journal



(b) Friendster



(c) PLD

Figure 5: Speedup over MPI nodes of different PageRank implementations on datasets from table 5

when the system is in a low thread and core count configuration. This performance can be due to a number of things such as:

- Message overhead vs growing nodes
- Parallelizability
- Computational overhead in addition to the baseline power iteration approach

The details of which will be addressed in a later part of this paper. This underperformance of the data-driven methods are not ideal as our project aims to propose methods that are faster than the topological Pagerank method. The implementation in a distributed system creates the goal that the methods are scalable in order to take advantage of the system.

On the other hand, the figures 2-5 show that not all cores are created equal. This refers to the time it takes for a Pagerank to run is not equal, even when there are the same number of cores but implemented in a different configuration. This is to say, X nodes by Y threads does not have the same runtime as Y nodes by X threads. The figures above exemplify that the implementations have lower run times when the number of MPI nodes grows. This is in contrast to the loss in speedup as additional threads are added in openMP. In many cases, the runtimes actually get slower on the 16 thread setting than the 1 thread setting. This could be attributed to the setup of the system. TACC, the computer system the computations were run on, composes 1 node as 2 Intel Xeon processors with 8 threads each. In the 16 thread case, the 16 cores are sharing the same pieces of memory across 2 processors. No matter where the data is stored, the processors have to access memory not in their cache, or maybe openMP has a method of automatically load balancing the memory, which would take time to compute what to do and move the memory.

The algorithm 1, topological implementation, iterates over all nodes it connects to in P_c , updates all of its \mathbf{x}_c , and then sends out an update for each of the nodes working on the problem. This breaks down into 1 iteration of P_c , sending a part of \mathbf{x} , and receiving a full \mathbf{x} . Although all x_v are updated in each node, the calculation is thread safe inherently.

The algorithm 2, pull implementation, iterates over all nodes, checks if they should be updated, then for each x_v that passes, it iterates over all P_v to calculate its new value. If x_v^{new} is high enough, it iterates over all of P_v^T to add its connections to the todo list. Then it sends the \mathbf{x}_c and merges its todo list with other nodes. The resulting overhead is one pass over P_c^T , and send and reduce a very long vector. This overhead might not be worth the fewer amount of \mathbf{x} 's computed.

Similarly, in Pull-push implementation, algorithm 3, the overhead is similar, but instead of a todo list, this method keeps a residual value. The update of this residual is not thread safe, as multiple nodes can be connected to and required to update the same node at the same time. To mitigate this, our implementation keeps separate thread copies of the residual thread and combines them after all of the \mathbf{x}_c^{new} are computed. Then, it is sent in the same fashion as Pull.

In Push implementation, algorithm 4, the method only needs one pass over P_c^T . This lessens the amount of memory accesses and computation needed for each iteration. This means that the overhead is the lowest of all of the methods. The performance is reflected in figures 1-5, as this is the only method that had consistently faster times than the baseline, Topological.

After considering the physical limitations of the methods, the results make sense.

6.2 ADMM Results Compared to other Linear methods

The results below are an attempt to compare other parallel linear solvers for $A\mathbf{x}=\mathbf{b}$ on different data sets and machines with our results for ADMM. This is not the primary focus of the ADMM results,

but it may offer some general intuition about the ADMM method compared to linear solvers not addressed in this paper. The residual value, after which the process stops, for all data below was set to 10^{-7} and the data set used for the ADMM method was the livejournal dataset, with ρ set to 10, and λ set to 4.

Method	Nodes	Iterations	Total Time (s)	Processes
Jacobi	18.5M	71	9.94	60
PR	2M	84	5.04	20
ADMM	1.7M	345	38.8	16
ADMM	1.7M	643	28.15	64

Again, although it is difficult to compare results on different data sets from different machines, the above results would indicate that the ADMM method (with gradient descent minimization) does not perform well compared to other parallel methods. This poor performance is likely due to a few things: As we increase the number of processes to solve the problem, the amount of information that has to be passed through MPI becomes a bigger burden and this ADMM code has not been optimized for message passing. Also, the minimization solver could likely be improved by changing the solver to something that is more appropriate given the underdetermined systems that are solved. Our ADMM solver, unlike our data-driven approach, does not take into consideration any information about the structure or values of our data. It would be possible to optimize according to some update criteria; this will be expanded upon in the future work section of this paper. Also, the software used in the competing methods (PETSC) has been significantly optimized and therefore we generally should expect our naive method to perform worse, regardless of the specific solver.

In addition to comparing the ADMM method to other parallel ($Ax=b$) solvers, the data set below will serve as a way to compare ADMM to the power iteration methods (push, pull, push-pull, topological). The DataSet column below will indicate the livejournal(LJ) or friendster(FR). We chose to not run the ADMM method on any larger data sets as the results for this method tended to be poor and the primary purpose of this paper is to focus on the results of distributing the data-driven methods.

DataSet	Iterations	Total Time (s)	Processes
LJ	345	38.8	16
LJ	643	28.15	64
LJ	1043	36.2	128
FR	420	287	16
FR	771	190	64
FR	1,295	201	128

The results above would indicate that there is the expected speed up when we increase the number of processes, but it is without question not a linear speedup. Again, we suspect this is due to the increase in message overhead. It is also indicated that the time per iteration increases which is exactly what we expected. Each iteration should take less time as the amount of data on each process will decrease with an increase in processes. Additionally, the number of iterations increases which we also expected since each individual process has an underdetermined process and will quickly solve for some minimum value but may not be close to the final global solution (dues to many local minimums) and therefore require much guidance (updates from the global variable z) to converge. Finally, we notice that with 128 processes, the total solve time begins to increase. We suspect this again due to inefficient message passing consuming a significant amount of time.

Although ADMM appeared to have poor results relative to other tested methods, it should be noted that the coding for ADMM was relatively simple and easy to understand. Aside from getting familiar with MPI and the minimization solver, there was nothing difficult about creating the code. A valuable metric for those looking for a quick (not so efficient) method for solving problems in this category.

7 Future Work

The data-driven algorithms update values that are possibly not thread safe. In the case of the residuals, updating that information created a large overhead in the multi-thread case. Using a framework like Galois that accomodates a lightweight atomic update, or updating a buffer to add onto an vector value would get rid of the need to keep separate thread copies of all of the residuals.

One big overhead was the vector of values that were checked every iteration. To lower this, and to lower the amount of information sent, we suggest another improvement would be to move to a worklist implementation of the data-driven methods. We did not implement a work list because the residuals were spread between each core, and therefore we could not implement a low overhead way of combining then creating a worklist.

Eventhough our method load balanced by nonzeros in each node, the nodes which continued to require recomputation were not necessarily equally distributed throughout the nodes. This meant that some nodes might not have had many values updated, but others would compute many. Because each node synchronized at the end of an iteration, the fastest iteration was as slow as the slowest node. To address that issue, there could be more dynamic scheduling of indices. But, as to be able to run on the TACC computers for the bigger datasets, the connection graph was distributed across nodes, and therefore, reassigning indices was not a viable option for the system.

Related to ADMM, there are many possible areas to focus for future research. Optimization can be done around selecting a better method for minimization on each process. Not only around the general method, but modifying that method to take a more data-driven approach since we have seen this can have a significant impact on results. Our approach uses the lasso method (L^1) regularization with the intent to create a sparse solution therefore reducing message overhead if using a sparse matrix and vector representation, which proved to be less efficient than standard solvers.

Aside from updating the minimization solver, the ADMM method could be improved by optimizing the message passing process. The current code does not take into consideration sparsity of data nor the change in values it updates; meaning if a values was being updated to a values that was not significantly different than its original value, can there be significant speedup with deciding not to update those values. This could be tremendously valuable when updating the global value z and broadcasting that update back to all processes.

8 Conclusion


In this paper, we overviewed the distributed Pagerank implementations of Topological, data-driven (pull, pull-push, and push), and ADMM. It compared the results over a varying amount of nodes and threads. In the end, the best implementations were push data-driven and topological Pagerank. Each of their performance was better than the other in some cases. Even though the data-driven methods have been shown to be consistently faster than topological implementations in Galois [2], artifacts of openMP and MPI disallowed the elements of their implementation that created the speedup. These places of potential improvement, outlined in Future Work, could be implemented

given enough time. Overall, the only method that outperformed the baseline semiconsistently was the Push data-driven Pagerank method. The push data-driven implementation outperformed the baseline on low openMP threads and MPI nodes, as well as large datasets. Low openMP thread and MPI nodes condition is not good for scalability, but the difference in performance over large datasets is promising.


This investigation also revealed that ADMM (without any optimization with respect to message passing nor the minimization solver) performed poorly relative to the other linear solvers reviewed in the this and other papers. In addition to the potential for better performance with optimization, ADMM can be a valuable technique for those wishing to use a method that is simple to code and easy to understand.

Appendices

Additional datasets' runtime, scalability, and speedup figures.




(a) SD1




(b) Twitter Rv

Figure 6: Convergence times of different PageRank implementations on larger datasets from table 5




(a) SD1




(b) Twitter Rv

Figure 7: Scalability over openMP threads of different PageRank implementations on datasets from table 5




(a) SD1




(b) Twitter Rv

Figure 8: Scalability over MPI nodes of different PageRank implementations on datasets from table 5




(a) SD1




(b) Twitter Rv

Figure 9: Speedup over openMP threads of different PageRank implementations on datasets from table 5



(a) SD1



(b) Twitter Rv

Figure 10: Speedup over MPI nodes of different PageRank implementations on datasets from table 5

References

- [1] Stephen Boyd, Neal Parikh, Eric Chu, Borja Peleato, and Jonathan Eckstein, “Distributed optimization and statistical learning via the alternating direction method of multipliers,” *Found. Trends Mach. Learn.*, vol. 3, no. 1, pp. 1–122, Jan. 2011.
- [2] Joyce Jiyoung Whang, Andrew Lenharth, and Inderjit S. Dhillon, “Scalable data-driven pagerank: Algorithms, system issues & lessons learned,” Submitted.
- [3] David Gleich, “Fast parallel pagerank: A linear system approach,” Tech. Rep., 2004.
- [4] Xiaoyi (Eric) Li, “Parallel pagerank computation using mpi,” <http://www.cse.buffalo.edu/faculty/miller/Courses/CSE633/Li-Fall-2012-CSE633.pdf>.
- [5] Asser N. Tantawi and Don Towsley, “Optimal static load balancing in distributed computer systems,” *J. ACM*, vol. 32, no. 2, pp. 445–465, Apr. 1985.