

Distributed PageRank - Final Report

Aaron Myers, Megan Ruthven

December 7, 2014

Contents

| | | |
|----------|---|----------|
| 1 | Introduction | 2 |
| 2 | Linear System Approach | 2 |
| 2.1 | ADMM | 3 |
| 2.2 | ADMM Results Compared to other Linear methods | 4 |
| 3 | Power Iteration Approach | 5 |
| 3.1 | Power Graph Results | 5 |
| 4 | Next Steps | 6 |
| 4.1 | Load Balancing | 6 |
| 5 | Related Work | 6 |
| 6 | Results | 6 |

1 Introduction

The purpose of this project is to investigate distributed PageRank by applying Pagerank methods currently in use in the multi-thread case to the distributed scenario and determine if these methods can be scaled to multi-machine systems. We also take a method that has already proven to be effective in the distributed setting and apply it to the Pagerank problem (ADMM [3]). Immediately below is a more detailed description of each approach. The primary metrics for performance will be: speedup, scalability, and ease of coding/understanding, all of which will be explained in the results section. The last metric is included to suggest that ease of coding and understanding of a method has an impact on the adoption rate in industry and therefore should be included in this investigation.

1. Apply ADMM [3] to the linear form of the Pagerank problem ($Ax=b$) and determine if there is value in taking this approach rather than other methods (Power Iteration, Linear solver with GMRES, BICGSTab, etc). Since it is more difficult to apply a data-drive approach to the separated minimization problems, we expect this method to have less than great performance for speed, but higher performance for ease of coding/understanding.
2. Approach the problem with the typical power iteration combined with data-driven and topology-driven approaches (listed below) to iterating along with appropriate load balancing. The idea is to maintain a work-list that contains only nodes whose change in pagerank value is above a set threshold, we will refer to this as "delta updating". This method is primarily taken from [4] and we will simply apply this method in the distributed setting with MPI. There are both major and subtle differences between each of the approaches below and we will go into greater detail about these differences later in this paper, although a more thorough explanation can be found in [4].
 - (a) Push based:
 - (b) Pull based:
 - (c) Push-pull based:
 - (d) Topology driven:

2 Linear System Approach

This approach requires that we form the PageRank problem into a linear system ($Ax=b$) in which case solving for x would provide the list of PageRank values. Below is a simple derivation taken from [2].

$$P' = P + dv^T \tag{1}$$

$$P'' = cP' + (1 - c)ev^T \tag{2}$$

$$x^{k+1} = P''^T x^k \tag{3}$$

Where P' and P'' are the modified PageRank matrices that have the modifications necessary to create a connected graph and add a personalization factor and e is a vector of all 1's, resulting in equation 3, the Power Iteration approach to Pagerank.

Given the additional information below, we can derive the linear system for Pagerank.

$$e^T x = x^T e = \|x\|_1 = \|x\| \quad (4)$$

$$d^T x = \|x\| - \|P^T x\| \quad (5)$$

$$x = [cP^T + c(vd^T) + (1-c)ve^T]x \quad (6)$$

Combining the information above, we arrive at the following equation:

$$(I - cP^T)x = kv \quad (7)$$

We now have Pagerank in a linear form ($Ax=b$), where $A = I - cP^T$ and $kv = b$. If in addition, we normalize x , we have the following:

$$k = \|x\| - c\|P^T x\| = (1-c)\|x\| + d^T x \quad (8)$$

$$k = 1 - c \quad (9)$$

2.1 ADMM

Many of the articles we encountered for solving parallel pagerank in the linear form used Jacobi iteration or some Krylov Subspace method (GMRES, BiCGSTAB, etc), but we attempted to implement something we were introduced to in this course, namely ADMM [3]. This is an extremely simple way to parallelize a linear solve. This process attempts to split the linear problem into subsections, solve separately, and combine the information in a very specific way. We will compare these results to the implementation of GMRES and BiCGSTAB for the same problem parallelizing using PETSc [1]. We expect ADMM to have worse performance, measured by speedup, but we would like to quantify the loss in accuracy/time relative to the ease of implementation and scalability.

Below is a brief description of the ADMM idea and algorithm [3].

We take the linear problem and split up the data accordingly:

$$A = [A_1 \dots A_n]' \quad (10)$$

$$b = [b_1 \dots b_n]' \quad (11)$$

Our original minimization of $Ax=b$ with a certain norm and regularization on x now becomes:

$$\text{minimize } \sum_{i=1}^N l_i(A_i x_i - b_i) + r(z) \quad (12)$$

$$\text{subject to } x_i - z = 0 \quad \forall i \quad (13)$$

Where x_i are local variables that we force to match the global solution z at each step and N is the number of processes used to solve the problem.

The resulting algorithm, using the augmented lagrangian presented in the ADMM method [3], is as follows:

Algorithm 1 ADMM Iteration

- 1: $x_i^{k+1} = \text{argmin}_x l_i(A_i x_i - b_i) + \frac{\rho}{2} \|x_i^k - z^k - u_i^k\|_2^2$
 - 2: $z^{k+1} = \text{argmin}_z r(z) + \frac{N\rho}{x} \|z^k - \bar{x}^{k+1} - \bar{u}^k\|_2^2$
 - 3: $u_i^{k+1} = u_i^k + x_i^{k+1} - z^{k+1}$
-

Where $u_i^k = \frac{1}{\rho} y_i^k$ and for our implementation, we chose the L^1 regularization term (also known as lasso) and used a gradient solver with Eigen to solve the minimization for x . Using the lasso regularization the z update becomes soft-thresholding update.

2.2 ADMM Results Compared to other Linear methods

Below are the initial results for ADMM programmed in Matlab for a simple data set (a disconnected synthetic 11 node graph).

Table 1: Linear Method Table

| 5 iter results | |
|----------------|-------------------|
| Method | $\ \hat{x} - x\ $ |
| GMRES | 0.0047 |
| BiCGSTAB | 0.0056 |
| ADMM | 0.0079 |

3 Power Iteration Approach

In addition to the Linear system, we will approach Distributed PageRank as a Graph with power iteration. We have slightly modified the approach (inspired by [4]) by first computing all updated pagerank values and for every subsequent update, we only modify the pagerank of the nodes pointed to by an update change in value above some certain threshold. We will also use the magnitude of these changes to prioritize a worklist for the algorithm to execute.

Algorithm 2 Power Iteration with Worklist

```

1: Initialize  $x, \delta$  (threshold)
2: Compute  $Px$  for all nodes
3: while Worklist is not empty do
4:   if  $x_i$  in worklist then
5:     take  $x_i$  off the worklist
6:      $x_i^{new} = (1 - \alpha) * P_i * x + \frac{\alpha}{\# [x]}$ 
7:     if  $|x_i^{new} - x_i| > \delta$  then
8:        $x_i = x_i^{new}$ 
9:       add  $x_j$  onto worklist :  $\forall x_i \rightarrow x_j$ 
10:    end if
11:  end if
12: end while

```

3.1 Power Graph Results

We implemented the two algorithms of pagerank with openMP, and compared the time to convergence. On the A dataset from HW4, the pagerank values converged in fewer iterations (25 to 20) and on average, takes less time to run one iteration (0.189 to 0.095 seconds) running on 16 threads. This resulted in a total calculation time of 4.72 seconds for the baseline power iterations and 1.90 seconds for the delta method. A summary table is below.

Table 2: Power Iteration - A

| Method Comparison | | |
|-------------------|-----------------|---------|
| Method | Iteration Count | Time(s) |
| Power Iteration | 25 | 4.72 |
| Delta Update | 20 | 1.90 |

Table 3: Power Iteration - Friendster

| Method Comparison | | |
|-------------------|-----------------|---------|
| Method | Iteration Count | Time(s) |
| Power Iteration | 23 | 49.6 |
| Delta Update | 30 | 19.5 |

4 Next Steps

Now that we have seen the value of both ADMM and data-driven PageRank, our next steps will involve parallelizing using MPI

1. Parallelize the ADMM method using MPI and compare the results to running GMRES and BiCGSTAB with PETSc
2. Parallelize the data-driven PageRank problem using MPI and compare these results to other parallel Power Iteration approaches
3. Collect all results and provide conclusions about all methods and the value that each provides.

4.1 Load Balancing

We will also implement load balancing with both ADMM and the delta updating as described below.

1. ADMM: each iteration, we will determine if the new local variable value has changed significantly. If so, it will be pushed to the master worklist to be operated on for the next iteration. If not, it will be removed from the worklist.
2. Delta Updating: similarly, each iteration will check the updated value of the node and push it and its out-neighbors to the master worklist if above a certain threshold. This master worklist will have duplicated removed and will divide the work evenly across all computing nodes. Ideally, the more connected nodes would be sent to the same compute node so some type of clustering may be beneficial for this operation.

5 Related Work

6 Results

References

- [1] David Gleich, et al. *Scalable Computing for Power Law Graphs: Experience with Parallel PageRank*
- [2] David Gleich, et al. *Fast Parallel PageRank: A Linear System Approach*
- [3] Stephen Boyd, et al. *Distributed Optimization and Statistical Learning via the Alternating Direction Method of Multipliers*
- [4] Joyce Whang, et al. *Scalable Data-driven PageRank: Algorithms, System Issues & Lessons Learned*
- [5] mcs.anl.gov <http://www.mcs.anl.gov/research/projects/mpi/tutorial/mpiexmpl/src/jacobiempl/C/main.h>
- [6] fsu.edu <http://people.sc.fsu.edu/~jburkardt/cppsrc/mpi/mpi.html>
- [7] Xiaoyi (Eric) Li *Parallel PageRank Computation using MPI*

- [8] David Gleich, Leonid Zhukov, Pavel Berkhin *Fast Parallel PageRank: Methods and Evaluations*
- [9] Shubhada Karavinkoppa and Jayesh Kawli *Page Rank Algorithm Using MPI*
- [10] Bundit Manaskasemsak, Putchong Uthayopas, Arnon Rungsawang *A Mixed MPI-Thread Approach for Parallel Page Ranking Computation*